

GEOG5415M Programming for Spatial Data Science

✓ Week 7: Machine Learning in Action

In this practical we will go through a worked example of machine learning in action. We are going to replicate some of the work in the paper:

- Asher, M., Oswald, Y., & Malleson, N. (2025). Understanding pedestrian dynamics using machine learning with real-time urban sensors. *Environment and Planning B: Urban Analytics and City Science*, 52(8), 1994-2017.
DOI:[10.1177/23998083251319058](https://doi.org/10.1177/23998083251319058)

The aim of the work is to build a model that can predict *pedestrian footfall* for a particular hour, given information about the time, the weather, and the built environment.

If you are interested, the original code, in full, is available from the paper's Github repository:

- github.com/nickmalleson/footfall/tree/main/MelbourneAnalysis

We will go through the following steps:

1. Data preparation (including downloading, cleaning and linking)
2. Data analysis (look for missing data and look at data distributions)
3. Model the data (including model selection)

Double-click (or enter) to edit

```
import os
import time

# Normal data science libraries
import geopandas as gpd
import pandas as pd
import seaborn as sns
#from scipy import stats
import numpy as np
import matplotlib.pyplot as plt

# machine learning packages
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
```

```
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV, K
from sklearn.metrics import mean_squared_error, r2_score

# set seaborn plotting theme to white
sns.set_theme(style="white")
```

```
# We need this to use the gdf.explore() function
!pip install mapclassify
```

```
Requirement already satisfied: mapclassify in /usr/local/lib/python3.1
Requirement already satisfied: networkx>=3.2 in /usr/local/lib/python3
Requirement already satisfied: numpy>=1.26 in /usr/local/lib/python3.1
Requirement already satisfied: pandas>=2.1 in /usr/local/lib/python3.1
Requirement already satisfied: scikit-learn>=1.4 in /usr/local/lib/pyt
Requirement already satisfied: scipy>=1.12 in /usr/local/lib/python3.1
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/li
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/d
```

✓ Data Preparation

We need to download, process, clean and merge the following data sources:

- footfall counts (our target)
- weather conditions
- built environment features
- date information (public holidays, school term times, etc,)

If you are interested, the full code is organised into various notebooks in the original repository's [PreparingData](#) directory.

✓ Downloading and reading data

✓ Footfall counts

The first dataset we need is the footfall counts -- i.e. the counts of pedestrians who pass Melbourne's footfall cameras every hour. In the paper, we had to find the data on the Melbourne Open Data Portal, download a few different files (one for old counts, one for new ones) and merge them. If you want to see the code to do this in full, have

a look at the [Data Preparation Folder](#) in the paper's repository. For this practical, we have made the data available with the notebook to save some time, but there is still a lot of cleaning and preparation to do.

The data are stored in the file called "sensor_counts.csv.gz" in the data/week_7 directory. Note that the file has the '.gz' extension. This is short for 'gzip' and it means that the csv file has been compressed using an algorithm called 'gzip' so that it takes up less space. Fortunately pandas makes it really easy to read and write files using the argument `compression='gzip'`. As with other data, we can read it directly from the module's github repository

Run the code below to read the compressed csv file and assign it to a variable called `sensor_counts`.

```
url_to_sensor_data = "https://github.com/MSc-Urban-Environmental-Leed:  
sensor_counts = pd.read_csv(url_to_sensor_data, compression='gzip')
```

Have a look at the sensor counts. Each row holds the number of counts from a sensor for a particular hour (the `hourly_counts`) column, as well as some other information. There are quite a lot of rows!

sensor_counts								
	Unnamed: 0	datetime	year	month	mdate	day	time	sen:
0	0	2019-11-01 17:00:00	2019	November	1	Friday	17	
1	1	2019-11-01 17:00:00	2019	November	1	Friday	17	
2	2	2019-11-01 17:00:00	2019	November	1	Friday	17	
3	3	2019-11-01 17:00:00	2019	November	1	Friday	17	
4	4	2019-11-01 17:00:00	2019	November	1	Friday	17	
...	...	...	...	...	...	...	...	...
5480637	1850385	2024-03-22 15:00:00	2024	March	22	Friday	15	
5480638	1850388	2023-05-27 10:00:00	2023	May	27	Saturday	13	

		13:00:00					
5480639	1850389	2024-03-23 16:00:00	2024	March	23	Saturday	16
5480640	1850390	2024-04-17 16:00:00	2024	April	17	Wednesday	16
5480641	1850391	2023-09-05 15:00:00	2023	September	5	Tuesday	15
5480642 rows x 9 columns							

There is a column that we don't need so let's get rid of it. The easiest way to do this is to use the `sensor_counts.drop()` function. For example, if you wanted to remove a column called 'MyColumn' you could remove it with:

```
sensor_counts.drop(columns = ['MyColumn'])
```

Write the code below remove the column called 'Unnamed: 0'. (Hint: don't forget to reassign the output of `drop()` back to your `sensor_counts` variable)

```
sensor_counts = sensor_counts.drop(columns = ['Unnamed: 0'])
```

The next chunk will check that the column removal worked correctly. If it runs without an error then you have done it right. (Ask someone if you'd like us to explain how it works)

```
column_still_in_df = [c for c in ['Unnamed: 0'] if c in sensor_counts]
if column_still_in_df:
    raise ValueError(f"These columns are still in the dataset: {column_still_in_df}")
print("Column removed successfully")
```

```
Column removed successfully
```

Now we need to check that the remaining columns that we have store their data correctly. Are they stored using an appropriate data type?

Edit the code below to see what types of data the columns store. (hint: you can use the `.info()` function to do this).

```
sensor_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5480642 entries, 0 to 5480641
```

```
Data columns (total 8 columns):
#      Column      Dtype
---  -
0      datetime    object
1      year         int64
2      month        object
3      mdate        int64
4      day          object
5      time         int64
6      sensor_id    int64
7      hourly_counts int64
dtypes: int64(5), object(3)
memory usage: 334.5+ MB
```

Most are OK. The months and days are text, so these are correctly represented as `object` types. The numerical data are all whole numbers, so correctly represented as `int`s.

But what about the `datetime` column? That is also being stored as a string object, which means that pandas won't be able to do reliable date arithmetic (e.g. comparing dates, adding to dates, etc.).

To fix this we can change the type of the column to a `datetime` object using the `pd.to_datetime()` function. Have a look at the [documentation](#). It tells us that the only 'positional' (required) argument is a Series (the `arg` parameter). Often we also need to provide the `format` argument to tell it how to convert the string to a date, but in this case our string date is formatted in a standard way so the function can work out what to do.

If we had a string column called `my_date` in the dataframe called `df` we could convert it like this:

```
df['my_date'] = pd.to_datetime['my_date']
```

Write the code below to convert the column called `datetime` column to a proper 'datetime' object

```
sensor_counts['datetime'] = pd.to_datetime(sensor_counts['datetime'])
```

This next chunk checks that the conversion worked ok

```
if not pd.api.types.is_datetime64_any_dtype(sensor_counts['datetime']):
    raise ValueError(f"The datetime column is incorrectly being stored")
else:
    print("✅ Datetime column is a datetime object")
```

✔ Datetime column is a datetime object

While we're doing time-related things, let's create columns to represent the month and day of week as a number. This may be more useful later than the text representations.

This is easy in pandas because, once we have created columns that have the type `datetime`, we can use `.dt` to access date- and time-related aspects. To find out what is available through the `.dt` attribute, have a look at the [dt docs](#) and, more usefully, follow the link to the [DatetimeIndex docs](#). Under 'attributes' on that page it shows you all the different parts of a date that you can access. For example, `sensor_counts['datetime'].dt.month` gives you the number of the month.

Edit the code below to fill in the variables. Then you will be able to print out the year, the day of the week, and the quarter for the very first row in our dataframe.

```
first_row = sensor_counts.loc[[0], :] # Use loc to access the first

year = first_row['datetime'].dt.year
# Do the same as above to create variables called 'day_of_week' and '
day_of_week = first_row['datetime'].dt.day_of_week
quarter = first_row['datetime'].dt.quarter

print("The date of the first row is:", first_row['datetime'].values)
print("The year, day of week and quarter are:", year.values, day_of_w

The date of the first row is: ['2019-11-01T17:00:00.000000000']
The year, day of week and quarter are: [2019] [4] [4]
```

Two things to note about the code above for those who are interested

1. Did you see that when we created `first_row`, we used an extra set of square brackets around the 0: `first_row = sensor_counts.loc[[0], :]`? If we didn't do this then pandas would realise that we only wanted one row, and instead of returning a DataFrame object it would return a Series. Series behave slightly differently and means the example wouldn't have worked. So adding the square brackets around the zero (`[0]`) means that we are giving `loc` a list, rather than a single number. This 'tricks' `loc` into returning a DataFrame (even though that DataFrame only contains one row).
2. In the `print` statement we use `.values`. This is because `.dt` returns a series, and when we `print` that we also get things like the row number, which

makes the print statement messy. `.values` means we *only* get the value, so it looks neater.

3. Question for those who really want to understand what's going on: why are our year, day and quarter values surrounded by square brackets when they are printed out?

Use `dt` to create columns to store the weekday and months as numbers

```
sensor_counts['weekday_num'] = sensor_counts['datetime'].dt.weekday + 1
sensor_counts['month_num'] = sensor_counts['datetime'].dt.month
```

It will also be useful later to have a datetime column that just has the day, not the full minutes and hours. This is so that we can join it to other datasets that have daily resolution (don't worry about this - ask someone if you're not sure).

```
sensor_counts['datetime_day'] = sensor_counts['datetime'].dt.floor('D')
```

There is one other date-related thing we need to do, related to **cyclical features**. Recall that in the lecture we discussed that times and days have big jumps (e.g. at midnight, from Sunday -> Monday and from December -> January). To get round this, we create *sine* and *cosine* transformations of those variables. Then the model is able to use two columns to capture each feature without the big jumps.

Run the code below to create new features for the time, month and weekday.

Don't worry about how this works, but if you would like to know please ask someone and we can explain it -- the concept isn't hard but the code includes some features that you haven't seen yet.

```
def add_sin_and_cos_features(df, column_to_transform):
    df['Sin_{}'.format(column_to_transform)] = np.sin(2 * np.pi * df[column_to_transform])
    df['Cos_{}'.format(column_to_transform)] = np.cos(2 * np.pi * df[column_to_transform])
    return df

for col in ['time', 'month_num', 'weekday_num']:
    sensor_counts = add_sin_and_cos_features(sensor_counts, col)
```

```
sensor_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5480642 entries, 0 to 5480641
Data columns (total 17 columns):
```

```
Data columns (total 17 columns):
```

```
#    Column          Dtype
---  -
0    datetime        datetime64[ns]
1    year            int64
2    month           object
3    mdate           int64
4    day            object
5    time           int64
6    sensor_id       int64
7    hourly_counts  int64
8    weekday_num     int64
9    month_num       int32
10   datetime_day    datetime64[ns]
11   Sin_time        float64
12   Cos_time        float64
13   Sin_month_num   float64
14   Cos_month_num   float64
15   Sin_weekday_num float64
16   Cos_weekday_num float64
```

```
dtypes: datetime64[ns](2), float64(6), int32(1), int64(6), object(2)
```

```
memory usage: 689.9+ MB
```

Start coding or generate with AI.

[] TODO Check for duplicates and other errors

✓ Sensor locations

We have now finished reading and preparing the footfall counts.

We know what the counts at each sensor are, but we don't know *where* the sensors are located. We need to know this so that we can map the sensor counts, and also so that we can join the sensors to other spatial data about the local built environment. The sensor locations are a separate file that we donloaded from the Melbourne Data Portal. For convenience we have put the file on module repository so you can load it easily:

```
sensor_locations = pd.read_csv("https://github.com/MSc-Urban-Environm
sensor_locations
```

```
      Unnamed: 0  sensor_id  Name  sensor_name  installation_date  stat
0              0          16  Australia  Col270_T  30/03/2009
```


1	1	50	Faraday St-Lygon St (West)	Lyg309_T	30/11/2017	
2	2	73	Bourke St - Spencer St (South)	Bou655_T	02/10/2020	
3	3	66	State Library - New	QVN_T	06/04/2020	
4	4	59	Building 80 RMIT	RMIT_T	13/02/2019	
...	...	...	...	...	...	
137	128	140	COM Pole 2837 - Boyd Park	Boyd2837_T	2023-11-02	N
138	129	150	narrm ngarrgu Library - Level 1 Main Stairs B	narrLibL1MB_T	2023-10-23	N
139	130	152	narrm ngarrgu Library - Level 2 - Study Area L...	narrLibL2S1_T	2023-10-23	N
140	131	154	narrm ngarrgu Library - Level 3 Children's Lib...	narrLibL3C1_T	2023-10-23	N
141	132	158	514 Elizbeth Street	Eli514_T	2024-01-29	N

142 rows x 13 columns

Next steps:

[Generate code with sensor_locations](#)[New interactive sheet](#)

Again, you need to get rid of that strange column.


Write the code below remove the column called 'Unnamed: 0'..

```
sensor_locations = sensor_locations.drop(columns = ['Unnamed: 0'])
```

Stretch activity: Convert the `sensor_locations` table to a GeoDataFrame and map it.

```
# Optional stretch activity

sensor_gdf = gpd.GeoDataFrame(
    sensor_locations,
    geometry=gpd.points_from_xy(sensor_locations['Longitude'],
                                sensor_locations['Latitude']),
    crs="EPSG:4326"    # WGS84 (lat/lon)
)
sensor_gdf.explore()
```



✓ Merge the sensor counts with the sensor locations

Now that we have counts for the sensors as well as their locations, we can merge the two DataFrames. This is made easy because both datasets have a column called `sensor_id`, which uniquely identifies a sensor.

We will use the `pd.merge` function to merge them. Have a quick look at the [pandas.merge\(\) documentation](#). The function has two *required* parameters: 'left' and 'right'. These are the two datasets that you want to merge. In our case it also needs the 'on' parameter so that it knows which column to use to merge the datasets. For example, if we wanted to merge two datasets, `df_a` and `df_b`, using column `id` as the linking column we would use `merge` like this:

```
merged_df = pd.merge(df_a, df_b, on='id')
```

Edit the code below to create a new variable called `location_counts` by merging the `sensor_locations` and `sensor_counts` dataframes using the `sensor_id` column.

```
location_counts = pd.merge(sensor_locations, sensor_counts, on='sensor_id')
location_counts
```

	sensor_id	Name	sensor_name	installation_date	status	
0	16	Australia on Collins	Col270_T	30/03/2009	R	n loc (22/C
1	16	Australia on Collins	Col270_T	30/03/2009	R	n loc (22/C
2	16	Australia on Collins	Col270_T	30/03/2009	R	n loc (22/C
3	16	Australia on Collins	Col270_T	30/03/2009	R	n loc (22/C

4	16	Australia on Collins	Col270_T	30/03/2009	R	n loc (22/0
...	...	...	...	...	...	...
5480637	140	COM Pole 2837 - Boyd Park	Boyd2837_T	2023-11-02	NaN	
5480638	140	COM Pole 2837 - Boyd Park	Boyd2837_T	2023-11-02	NaN	
5480639	140	COM Pole 2837 - Boyd Park	Boyd2837_T	2023-11-02	NaN	
5480640	140	COM Pole 2837 - Boyd Park	Boyd2837_T	2023-11-02	NaN	
5480641	140	COM Pole 2837 - Boyd Park	Boyd2837_T	2023-11-02	NaN	

5480642 rows × 28 columns

To make sure this worked, lets check that your new `location_counts` dataframe has the same number of rows as the original `sensor_counts` dataframe

```
if len(location_counts) == len(sensor_counts):
    print("Merged dataframe has the correct number of rows")
else:
    print("Merged dataframe length", len(location_counts), "doesn't mi
```

Merged dataframe has the correct number of rows

✓ Weather and holiday data

We need information about when school and public holidays take place, and also information about the weather. These were collected from a few places, see [the paper](#) if you're interested.

One thing to note about the code below is that we use the

`parse_dates=['date']` argument. This tries to automatically create a datetime object when reading the 'date' column, rather than storing it as a string. It means that we don't need to convert the column later (like we had to do for the sensor data). The `dayfirst` argument tells pandas that for two of our files the day comes first, not the year. (We could have used this feature when reading the sensor data earlier, but it was good to see how you have to do it manually).

✓ Load weather and holiday data

```
url_to_data = "https://github.com/MSc-Urban-Environmental-Leeds/GE0G5"

daily_rainfall = pd.read_csv(url_to_data+"DailyRainfallData.csv", par
public_holidays = pd.read_csv(url_to_data+"publicholidays.csv", parse
school_holidays = pd.read_csv(url_to_data+"schoolholidays.csv", parse
```

Edit the code below to have a look at those four dataframes that you just loaded. Do they *look* OK? Are the data stored in the right format? Are there any duplicate rows? Etc..

```
# Check that the data are OK in this cell. If you want to add new cel
```

Start coding or generate with AI.

Start coding or generate with AI.

✓ Merge the weather and holiday data with the sensor data

We want to attach these new dataframes to our sensor data. All of our dataframes have a column called `date`, so we can merge by comparing that column to our `datetime` day

```
datetime_day).
```

We can do this using the `left_on` argument, which tells `merge()` that column in the left dataframe that we want to join, and the `right_on` argument for the right dataframe. So in our case we want `left_on="datetime_day", right_on="date"`. (We didn't have to do this earlier because both dataframes had a column called `datetime` that we used to merge).

IMPORTANT! Another thing to note is that there may be some specific dates missing from our weather and holiday data. The default merge behaviour is to drop rows where there is no match. We can tell pandas not to do this by making it a 'left' join. That means keep all the rows in the 'left' dataset, and if it cannot find a matching value then insert 'na' into that row. We can use the `how="left"` argument to `merge()` and make sure that `location_counts` is the first dataset we pass (that one is treated as the 'left' dataset, the other is treated as the 'right' one).

Run the chunk below to see how to attach the daily_rainfall data to the sensor data.

```
# Do the merge
location_counts = pd.merge(location_counts, daily_rainfall, left_on="
# Drop the 'date' column that was part of 'daily_rainfall'. If we don
location_counts = location_counts.drop(columns='date')
```

Now edit the chunks below to merge the other three datasets to the sensors data. (Don't forget to drop the 'date' column after each merge).

```
# Merge public_holidays
```

```
# Merge school_holidays
```

```
# REMOVE FROM STUDENT COPY
for df in [public_holidays, school_holidays]:
    location_counts = pd.merge(location_counts, df, left_on="datetime
location_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5481050 entries, 0 to 5481049
```

Data columns (total 31 columns):

#	Column	Dtype
0	sensor_id	int64
1	Name	object
2	sensor_name	object
3	installation_date	object
4	status	object
5	note	object
6	Latitude	float64
7	Longitude	float64
8	location	object
9	Note	object
10	Location_Type	object
11	Status	object
12	datetime	datetime64[ns]
13	year	int64
14	month	object
15	mdate	int64
16	day	object
17	time	int64
18	hourly_counts	int64
19	weekday_num	int64
20	month_num	int32
21	datetime_day	datetime64[ns]
22	Sin_time	float64
23	Cos_time	float64
24	Sin_month_num	float64
25	Cos_month_num	float64
26	Sin_weekday_num	float64
27	Cos_weekday_num	float64
28	rainfall_mm	float64
29	public_holiday	float64
30	school_holiday	float64

dtypes: datetime64[ns](2), float64(11), int32(1), int64(6), object(11)
memory usage: 1.2+ GB

Now use `.info()` to check that the columns in the new `location_counts` dataframe look OK. Do you have the right columns and are they of the expected type?

```
# Check the data have the columns we're expecting:
location_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5481050 entries, 0 to 5481049
Data columns (total 31 columns):
#   Column              Dtype
---  -
0   sensor_id           int64
1   Name                object
2   sensor_name         object
3   installation_date   object
4   status              object
```

```

5  note                object
6  Latitude            float64
7  Longitude           float64
8  location            object
9  Note                object
10 Location_Type       object
11 Status              object
12 datetime            datetime64[ns]
13 year                int64
14 month               object
15 mdate               int64
16 day                 object
17 time                int64
18 hourly_counts       int64
19 weekday_num         int64
20 month_num           int32
21 datetime_day        datetime64[ns]
22 Sin_time            float64
23 Cos_time            float64
24 Sin_month_num       float64
25 Cos_month_num       float64
26 Sin_weekday_num     float64
27 Cos_weekday_num     float64
28 rainfall_mm         float64
29 public_holiday      float64
30 school_holiday      float64
dtypes: datetime64[ns](2), float64(11), int32(1), int64(6), object(11)
memory usage: 1.2+ GB

```

One other thing we need to do. Our public and school holiday input data only contained dates that are holidays. Note those dates that are not holidays. So after the merge, all the days that are holidays will have a value of 1, and those that aren't holidays will have 'na'. We can convert those NAs to 0 to show that it isn't missing data, it's a day that isn't a holiday:

```

location_counts['public_holiday'] = location_counts['public_holiday']
location_counts['school_holiday'] = location_counts['school_holiday']

```

P.S. Before moving on, here's a hint with merging data: Once we ran the code

`location_counts = pd.merge(...)` we created a new `locations_counts` dataframe and deleted the old one. This means that if we had done something wrong with the merge, we would have to go back and re-read the sensor data, re-merge it to the locations, etc. It's not a problem, but it can be annoying, especially if some of the earlier merge operations took a long time. So, if you want to check whether the merge looks OK before doing it, you can just call the `pd.merge()` function but not assign the resulting new dataframe to the `location_counts` variable. For example you could do this:

could do this.

```
temporary_df = pd.merge(location_counts, daily_rainfall, on="datetime")
```

and then analyse `temporary_df` to see if the merge worked first. Or, if you just do:

```
pd.merge(location_counts, daily_rainfall, on="datetime")
```

pandas will helpfully show you what the result of the merge will look like (although it doesn't save the result, so you can't do any further analysis)

✓ Check the merges worked OK - do we have missing data?

Just because the columns exist in the merged dataframe doesn't mean that the merge actually worked correctly. There may be some rows where pandas couldn't find a match.

A nice easy way to check whether the merge was OK is to see whether any 'na' values were introduced after the merge. This happens when there are some rows in one dataset that pandas can't find a partner for in the other data. We've seen the combination of `isna()` and `sum()` functions before.

Run the code below. Are there any rows with missing data?

```
location_counts.isna().sum()
```

	0
sensor_id	0
Name	0
sensor_name	0
installation_date	0
status	173639
note	4817812
Latitude	0
Longitude	0
location	0
Note	5473529
Location_Type	5307411

Status	5307411
datetime	0
year	0
month	0
mdate	0
day	0
time	0
hourly_counts	0
weekday_num	0
month_num	0
datetime_day	0
Sin_time	0
Cos_time	0
Sin_month_num	0
Cos_month_num	0
Sin_weekday_num	0
Cos_weekday_num	0
rainfall_mm	1205948
public_holiday	0
school_holiday	0

dtype: int64

The missing rainfall data are a problem. In a real setting we would need to explore this in detail. We could start with making a table that just contained the missing rows and looking for patterns, e.g.:

```
location_counts.l[location_counts['rainfall_mm'].isna(), ]
```

and then we could try to impute the rainfall. However for now we are just going to make the missing values 0.0 (not good practice!!)

```
location_counts['rainfall_mm'] = location_counts['rainfall_mm'].fillna(0.0)
```

```
# Now there are no missing values for the columns that matter  
location_counts.isna().sum()
```

	0
sensor_id	0
Name	0
sensor_name	0
installation_date	0
status	173639
note	4817812
Latitude	0
Longitude	0
location	0
Note	5473529
Location_Type	5307411
Status	5307411
datetime	0
year	0
month	0
mdate	0
day	0
time	0
hourly_counts	0
weekday_num	0
month_num	0
datetime_day	0
Sin_time	0
Cos_time	0
Sin_month_num	0
Cos_month_num	0
Sin_weekday_num	0

```

Cos_weekday_num      0
    rainfall_mm      0
    public_holiday   0
    school_holiday   0

```

```
dtype: int64
```

✓ Load and merge the built-environment spatial data

Now we want to join in other data that may influence pedestrian footfall. Things like hospitals, universities, train stations, etc., etc.

Preparing the spatial data quite laborious. We downloaded OpenStreetMap data and then used the spatial analysis functions available in geopandas to find the number of different types of objects (like trees, stations, bike stands, bus stops, etc.) within a certain radius of each sensor. If you are interested in how we did this have a look at the [FindFeaturesInProximityToSensors.ipynb](#) script in the paper's repository.

For this practical, we will just load the pre-processed data and join it to the sensors. The data that we will read is a table that has, for each sensor, the numbers of different types of physical features within a certain distance (400m in our case)

```

num_features_near_sensors = pd.read_csv("https://github.com/MSc-Urban-
num_features_near_sensors

```

	index	lights	street_inf	bikes	landmarks	memorials	trees	trans
0	1	574	1527	2	18	21	454	
1	2	551	1557	2	16	20	454	
2	3	498	1331	3	16	10	455	
3	4	373	1640	2	17	20	416	
4	5	339	1112	2	18	11	706	
...	...	...	...	...	...	...	...	
137	159	58	965	2	2	7	751	
138	160	113	928	2	9	8	650	
139	161	677	1014	2	15	13	1248	
140	162	60	1005	2	5	7	688	

141	163	669	587	2	3	2	675
-----	-----	-----	-----	---	---	---	-----

142 rows x 36 columns

Write the code below to merge the `location_counts` data frame with the `num_features_near_sensors` dataframe, using the `sensor_id` column that is common to both dataframes..

```
location_counts = pd.merge(location_counts, num_features_near_sensors
```

```
# This checks you did the merge correctly by checking that the 'trees'
# 'trees' is just one of the columns from the num_features_near_senso
if not "trees" in location_counts.columns:
    raise ValueError("There is no 'trees' column after merging, sometl
else:
    print("There is a 'trees' column after the merge")

# If we get here then no error was raised so presumably the trees colu
if location_counts['trees'].isna().all():
    raise ValueError("'trees' column exists but all values are NA - tl
else:
    print("Merge successful: 'trees' column found and contains valid "
```

```
There is a 'trees' column after the merge
Merge successful: 'trees' column found and contains valid data.
```

✓ Temporary Data sampling...

We have a count of footfall for every hour, on every day, over more than 10 years. That is a **lot** of data. We can process it on colab, but it takes a long time (anywhere from a few minutes to a few hours depending on what we need to do). That would make the rest of the practical very tedious. So here we are just going to extract data for the year 2015 (chosen arbitrarily). If you are running this on a fairly decent laptop, or don't mind waiting for a bit, then comment out the lines below and calculate models for the full data set.

```
#X = location_counts.loc[ location_counts['datetime'].dt.year == 2015
#y = location_counts.loc[ location_counts['datetime'].dt.year == 2015
location_counts = location_counts.loc[ location_counts['datetime'].dt
```


✓ Data Analysis

Before modelling, we need to analyse our data. Using a combination of statistical analysis and visualisations (using the code you wrote in [Week 4](#)), write some code below to get a better understanding of your data. You can be inventive and explore the different variables in any way that you want to, but if you're not sure what to do then you can follow the prompts.

Plot a histogram of the hourly counts.

```
# Histogram for hourly counts  
sns.histplot(location_counts['hourly_counts'])
```

<Axes: xlabel='hourly_counts', ylabel='Count'>



Plot histograms for three additional variables. Are these all OK?

```
# First histogram
```

```
# Second histogram
```


```
# Third histogram
```

Do a joint scatter plot for the hourly counts and the rainfall (Hint: you can use `sns.jointplot()` with `kind='scatter'`).

```
# Code for a scatter (rainfall and hourly counts)

sns.jointplot(x='rainfall_mm', y='hourly_counts', kind='scatter', da

<seaborn.axisgrid.JointGrid at 0x7f6273b79d90>
```




rainfall_mm

Well that isn't very clear. Try a kde plot instead..

```
# Code for a kde plot
```

```
sns.jointplot(x='rainfall_mm', y='hourly_counts', kind='kde', fill=True)
```

```
<seaborn.axisgrid.JointGrid at 0x7f629917cb60>
```



Do box plots to compare the variables: bikes; landmarks; memorials; transport_stops. Hint: an easy way to do this is select the columns that we want using `loc` and pass that to the `sns.boxplot()` function:


```
sns.boxplot(data=location_counts.loc[:, ['var1', 'var2', 'var3', 'var4']])
```



```
# Code for four boxplots
```

```
sns.boxplot(data=location_counts.loc[:, ['bikes', 'landmarks', 'memor.
```

```
<Axes: >
```




Finally, this plot adds up the total footfall over the course of a day and plots the daily footfall for each sensor for the month of May. It's messy but you can start to see some of the trends emerge

```
# Sum hourly counts per day per sensor
daily = location_counts.loc[location_counts.month_num==5,:].groupby(['

# Line plot: one line per sensor
sns.lineplot(
    data=daily,
    x='datetime_day',
    y='hourly_counts',
    hue='sensor_id',
    palette='tab10',          # qualitative colour scheme
    linewidth=1,              # thinner lines
)
```


<Axes: xlabel='datetime_day', ylabel='hourly_counts'>



✓ Modelling

Now we are ready to do some modelling. Recall from the lecture that we need to do three things:

1. **Model Selection** to choose the best *type* of model
2. **Hyperparameter tuning** to find the best parameters for our chosen type of model
3. **Model evaluation** to see how well our optimised model works on data that it hasn't seen before.



✓ Data Preparation

Before continuing, let's get our X (predictors) and y (target) datasets ready. This mostly consists of just choosing the rows and columns in the `location_counts` data frame that we want to use for modelling.

Let's start by selecting the rows that we want. Remember that each row holds the footfall count for a particular hour in our study period. During and after COVID, footfall patterns changed dramatically, so it is unlikely that a model will be able to predict footfall under *normal* conditions as well as during the COVID period. Therefore we will remove the COVID period from our training data.

Also remember that there is a convention to call our data X and y , so we will do that

Start with the y data (our target). Use `loc` to filter out all rows that were before 2020 (the rows) and select only the 'hourly_counts' column.

As a reminder, the syntax of `loc` is like this:

```
df.loc[ <rows_to_filter>, <columns_to_select> ]
```

```
y = location_counts.loc[ location_counts['datetime'] < '2020', 'hourly'
```

Now do the same for the X dataset (our predictors). This is slightly more complicated though because rather than just selection one column (the 'hourly_counts') we want to select all the columns that we think might be related to footfall. Lets define the columns we want in their own list to make things easier. There is no 'correct' list here. If you're interested you could come back and add or remove variables from this list and see how it affects the model (remember that `location_counts.columns` will show you all the columns that are available).

```
# Choose the columns we want to use as our predictor variables, making
predictor_columns = [
    # The time-related variables:
    'Sin_time', 'Cos_time',
    'Sin_month_num', 'Cos_month_num',
    'Sin_weekday_num', 'Cos_weekday_num',
    # Weather-related:
    'rainfall_mm',
    # Holidays:
    'public_holiday', 'school_holiday',
    # Built-environment
    'lights', 'street_inf', 'bikes', 'landmarks', 'memorials', 'trees',
    'transport_stops', 'bus-stops', 'tram-stops', 'metro-stations', 'big-car-parks',
    'buildings_2019', 'avg_n_floors_2019'
]
```

Now we can use `loc` like we did before.

Complete the code below to use the `loc` accessor to choose all rows before 2020 and the columns we defined in the `predictor_columns` list.

```
# Create the 'X' variable with rows before 2020 and columns from the
X = location_counts.loc[ location_counts['datetime'] < '2020', predic
```

Have a look at the two new dataframes (X and y). Do they have the right columns? Do they have the same number of rows?

Do a couple of checks below to see that X and y are as you expect.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

```
# This just checks that the number of rows are the same
if not len(X) == len(y):
    raise ValueError("The number of rows in X and y are different")
else:
    print("X and y have the same number of rows, which is what we want")

X and y have the same number of rows, which is what we want
```

✓ Model selection

Sometimes you know which kind of model will best suite your application, but other times you need to try a few different ones. In the full paper we tested three types of model:

1. **Linear regression:** our data are non-linear, so we don't expect this to work well, but it is a good 'benchmark'
2. **Random forest regression:** these handle non-linearity well
3. **XGBoost::** a gradient boosting method that can out-perform random forest

Here, to simplify the code slightly, we will just test **Linear regression** and **Random forest regression**.

Start by creating our scikit-learn `Pipeline` objects. Note: as part of the pipeline we can also use the `StandardScaler` to scale our data using Z-Scores.

```
# Pipeline for linear regression
lr_model_pipeline = Pipeline(steps=[
    ['scaler', StandardScaler()],          # Step 1: normalise ('sc
    ['linear_regressor', LinearRegression()] # Step 2: run the linear
])

# Pipeline for random forest
rf_model_pipeline = Pipeline(steps=[
    ['scaler', StandardScaler()],
    ['rf_regressor', RandomForestRegressor(n_jobs = os.cpu_count())]
```

1)

We could do full cross-validation with model selection, but for this example we won't. Instead, we will just create our own train and validate (test) datasets, using the handy `train_test_split` function that scikit-learn provides. We need four datasets: train/test for our X data, and train/test for our y data:

- `X_train`: a subset of the predictor data (weather, built environment, etc.) that we use to train the model
- `y_train`: a subset of the real football counts data that we use when training the model (the model will use the `X_train` data to try to make predictions that match these `y_train` data).
- `X_test`: a subset of the predictor data we use to *test* the model afterwards
- `y_test`: a subset of the real data that we use to test how well the model worked after we have finished training it.


Run the code below to have scikit-learn create test and train datasets.

```
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

Now we try fitting the linear regression and random forest models to those data (fit it to the training dataset and then test its accuracy on the test data)

Run the code below to fit the linear regression to the training data.


```
lr_model_pipeline.fit(X_train, y_train)
```



Now do the same for the random forest pipeline.

Write the code below to fit the linear regression to the training data.

```
# Fit the random forest just like the linear regression one above  
rf_model_pipeline.fit(X_train, y_train)
```



We have trained our models. Lets see which one worked best by seeing how well it can predict footfall, using the data in `X_test`, which it hasn't seen before, to predict footfall in `y_test`.

We will start with the linear regression:

```
# Create a new dataset of model predictions, giving the model the X_test
lr_predict = lr_model_pipeline.predict(X_test)

# Now calculate two error metrics to see how well it did
lr_rmse = np.sqrt(mean_squared_error(y_test, lr_predict))
lr_r2 = r2_score(y_test, lr_predict)
print("The linear regression model got rmse and r2 scores of:", lr_rmse,
```

The linear regression model got rmse and r2 scores of: 634.58630397948

OK, so we have used the fitted linear regression model to make predictions and calculated how good those predictions were when compared to some hourly counts.

Fill in the code below to make predictions for the random forest model and calculate how well the model worked.

```
# Create a new dataset of model predictions, giving the model the X_test
rf_predict = rf_model_pipeline.predict(X_test)

# Now calculate two error metrics to see how well it did
rf_rmse = np.sqrt(mean_squared_error(y_test, rf_predict))
rf_r2 = r2_score(y_test, rf_predict)
print("The random forest model got rmse and r2 scores of:", rf_rmse,
```

The random forest model got rmse and r2 scores of: 214.43296858880598

Which model was better??

Fill in the markdown cell below explaining which model performed better and how you know this.

...

We can also have a look at how the footfall predictions compared to the real footfall counts, to see if there are any systematic under- or over-predictions. If the predictions are perfect, then all points should sit on the $y=x$ line. This isn't strictly necessary at this stage but we will do more of this kind of analysis after optimising the final model.


Read the comments and run the code below, we will return to this later. Does this support your assessment of which model is better?


```
# The graphing is easier if we have all of our data (test data and the
# We can build the dataframe by passing it a dictionary object that we
temp_df = pd.DataFrame({
    'actual': y_test, # The test data ('real' values)
    'linear_regression': lr_predict, # Linear regression predictions
    'random_forest': rf_predict # Random forest predictions
})
```

```
lr_graph = sns.jointplot(data=temp_df, x='actual', y='linear_regression')
lr_graph.ax_joint.axline((0, 0), slope=1, linestyle='--', linewidth=1)

rf_graph = sns.jointplot(data=temp_df, x='actual', y='random_forest',
rf_graph.ax_joint.axline((0, 0), slope=1, linestyle='--', linewidth=1)
```

<matplotlib.lines.AxLine at 0x7f623c123440>





✓ Hyperparameter Tuning

The random forest model performed much better than the linear regression model, so we will use that for the rest of the analysis.

Now that we have selected our model, we can tune the hyperparameters to see if we can make it fit even better. There are a lot of hyperparameters that will influence how well a model works on a particular dataset. For example, the [RandomForestClassifier documentation](#) lists a large number of parameters that might influence how well it works. Here are just two:

- `n_estimators`: the number of trees in the forest (default 100). More trees can increase accuracy but can make it take longer to fit and use more computer memory.
- `max_depth`: the maximum depth of each tree (default None). Smaller trees are simpler and less susceptible to over-fitting. The default, *None* allows them to grow fully and risks overfitting.

To add these to our pipeline we need to prepare a dictionary in the following format:

```
{'model_name__parameter_name': [ list_of_values ] }
```

We called our model 'rf_regressor' when we created the pipeline, so we can specify parameters as follows:

parameters as follows:

```
parameters = {
    'rf_regressor__n_estimators': [50, 100],
    'rf_regressor__max_depth': [None, 10],
}
```

Edit the code below to add one additional parameter of your choice to the test

```
parameters = {
    'rf_regressor__n_estimators': [50, 100],
    'rf_regressor__max_depth': [None, 10],
    #'rf_regressor__min_samples_leaf': [1, 3, 5],
    #'rf_regressor__max_features': ['sqrt', 'log2']
}
```

There are a few different methods that the algorithm can use to try out the different parameters. Here we use [Grid Search Cross Validation](#). This basically tries all possible parameter combinations. If we have loads of parameters we might try a cleverer search algorithm such as the [Bayes Search CV](#) that tries guess which parameter combinations will work well and not bother testing those that wont.

Read the code below and run it to prepare the Grid Search Cross Validation.

```
grid = GridSearchCV(
    estimator=rf_model_pipeline,          # The pipeline we created earl
    param_grid=parameters,                # Our parameters
    scoring='r2',                          # How to calculate error ('neg
    cv=KFold(n_splits=2, shuffle=True),    # We want to do K-old cross va
    n_jobs=os.cpu_count(),                # If possible, run the job ovr
    verbose=2                             # Ask it to print some messag
)
```

Now we are ready to go. We call

```
grid.fit(X_train, y_train)
```

to run the grid search. Just for interest we also record the start and end time so we can see how long it takes to run.

Run the code below and go and make a cup of tea while you're waiting. When I ran this on colab it took about 10 mins, but this may vary depending on the number of parameter combinations you have.

```
start = time.time()
grid.fit(X_train, y_train)
end = time.time()

print("Grid search finished in", (end - start)/60, "minutes.")
```

```
Fitting 2 folds for each of 4 candidates, totalling 8 fits
Grid search finished in 7.8273456732432045 minutes.
```

Phew. It's done. Let's see which parameters were best:

```
print("Best params:", grid.best_params_)
```

```
Best params: {'rf_regressor__max_depth': None, 'rf_regressor__n_estima
```


Lets rebuild the model using those parameters and fit it to the data

✓ Model Evaluation

We have now selected our model and tuned its hyperparameters. The final thing is to see how well it can predict football on data it has not seen before.

Begin by creating a new model with those optimal hyperparameters. This is our 'best' model:

```
best_model = rf_model_pipeline.set_params(**grid.best_params_)
best_model.fit(X_train, y_train)
```



Now we can use that model to make predictions on our *test* data. It has not seen these data before. We evaluate it in the same way that we evaluated the models during model selection:

```
# Use the model to make predictions from our test data
best_model_predictions = best_model.predict(X_test)
```

```
# Calculate the error by comparing our test data to the model predict.
rmse = np.sqrt(mean_squared_error(y_test, best_model_predictions))
r2 = r2_score(y_test, best_model_predictions)

print("Test RMSE:", rmse)
print("Test R^2:", r2)
```


```
Test RMSE: 213.98861843815214
Test R^2: 0.9399156520029938
```

We can also graph the predictions against the real data as we did before

```
temp_df = pd.DataFrame({
    'actual footfall': y_test, # The test data ('real' values)
    'best model predictions': best_model_predictions # Random forest |
})

best_graph = sns.jointplot(data=temp_df, x='actual footfall', y='best
best_graph.ax_joint.axline((0, 0), slope=1, linestyle='--', linewidth:

<matplotlib.lines.AxLine at 0x7f623ca43ce0>
```



✓ Using the model to predict footfall

Now that we have a working model, there are loads of things we can do with it. For example:

- Analyse the errors to see where footfall varied from what we would have expected and try to understand why
- Plug in new X values to predict what footfall might be like next week
- Quantify how successful previous events were (did we see an increase or decrease in footfall given the weather conditions on the day?)

To finish, let's look at footfall on 1st January 2015. We will make a prediction of footfall at every hour on that day for one sensor in particular (sensor 5) and see how it compared to real footfall.

The code below should just run. Read it and please ask if you don't understand any of it. It looks complicated but there is nothing you haven't come across before.

```
# Create X and y variables for the 1st Jan 2019
X = location_counts.loc[
    (location_counts['datetime'].dt.date == pd.to_datetime('2015-01-01')) &
    ~location_counts['datetime'].dt.isna()
][predictor_columns]


y = location_counts.loc[
    (location_counts['datetime'].dt.date == pd.to_datetime('2015-01-01')) &
    ~location_counts['datetime'].dt.isna()
]['hourly_counts']
```


```
# Predict footfall on that date
prediction = best_model.predict(X)
```

```
plt.figure(figsize=(8,4))

sns.lineplot(x=range(len(y)), y=y, label='Actual')
sns.lineplot(x=range(len(y)), y=prediction, label='Prediction')

plt.xlabel("Time of day (hours)")
plt.ylabel("Hourly counts")
plt.title("Actual vs Prediction")
plt.tight_layout()
```





We have no mechanism in the model to tell it to expect very high footfall over night on 1st January, so it isn't surprising that our prediction is much lower over night than we saw in reality.

That's it! For those who are really keen:

Stretch activity: create your own hypothetical X data for a single sensor and hour, and make a prediction on what the footfall might be.